

07 – NoSQL

NoSQL, an alternative to SQL databases

Jérôme Landré – jerome.landre@ju.se

(Several slides have been created by Jasmin Jakupovic)

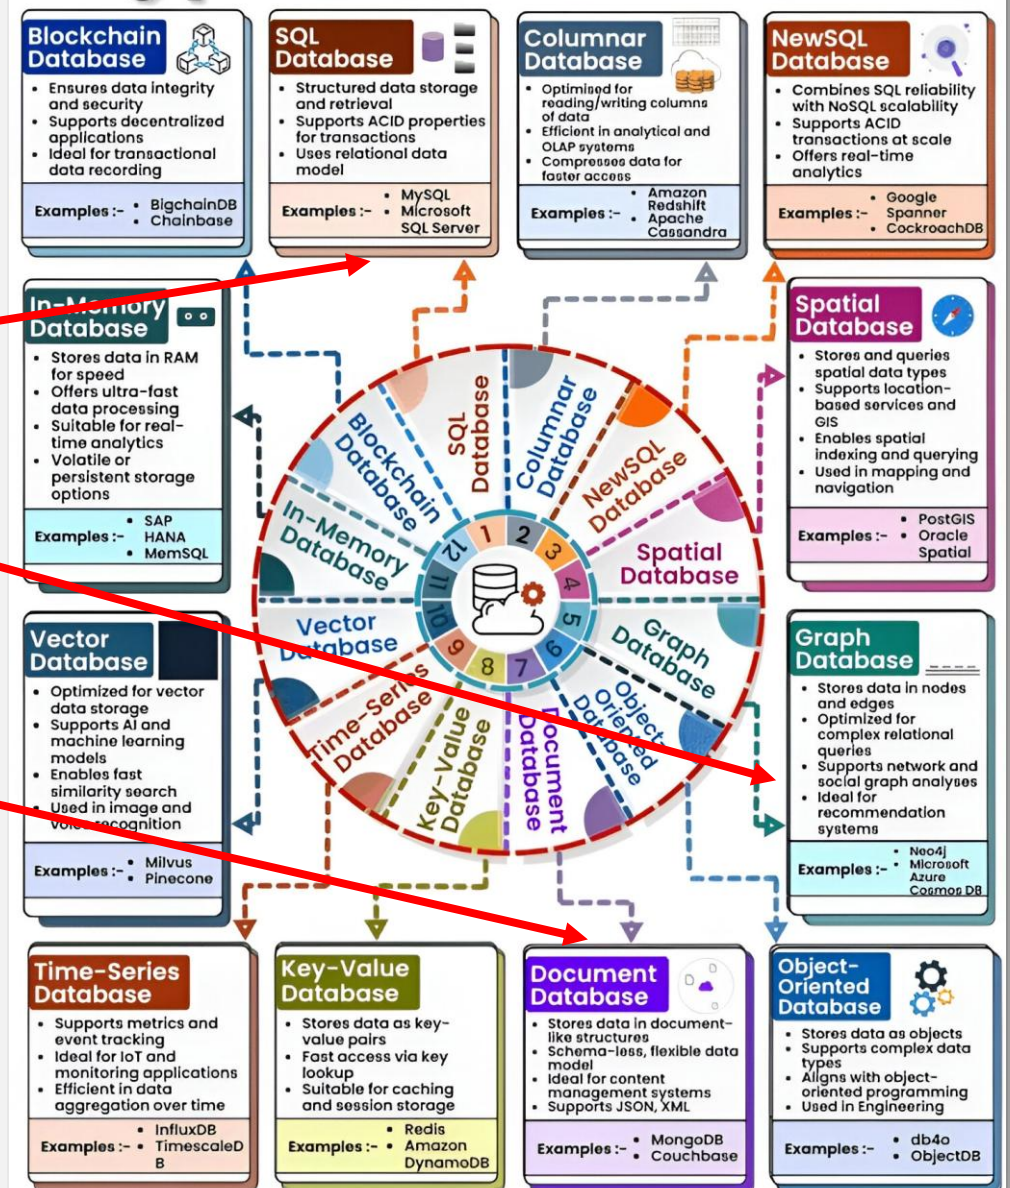
OUTLINE

- 1) Types of Databases
- 2) Databases rankings
- 3) NoSQL databases
- 4) Document databases
- 5) Graph Databases
- 6) Neo4J
- 7) Cypher

1) TYPES OF DATABASES

- There are many different types of databases, they can be sorted by type of data / type of application
- SQL Database, Graph Database, Document Database, Key-Value Database...

Types of Databases



2) DATABASES RANKING

- Relational DB are the most represented and the best ranked
- MongoDB, Databricks, Redis, Apache Cassandra are the four NoSQL DB in positions 5, 7, 8 and 10 respectively

Source: <https://db-engines.com/en/ranking>

431 systems in ranking, April 2026

Rank			DBMS	Database Model	Score		
Apr 2026	Mar 2026	Apr 2025			Apr 2026	Mar 2026	Apr 2025
1.	1.	1.	Oracle	Relational, Multi-model ⓘ	1157.93	-24.53	-73.12
2.	2.	2.	MySQL	Relational, Multi-model ⓘ	857.69	-0.65	-129.41
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model ⓘ	702.08	-9.40	-82.93
4.	4.	4.	PostgreSQL	Relational, Multi-model ⓘ	681.35	+1.27	+14.10
5.	5.	5.	MongoDB	Document, Multi-model ⓘ	385.02	+1.44	-15.03
6.	6.	6.	Snowflake	Relational	217.38	+6.13	+49.28
7.	7.	↑12.	Databricks	Multi-model ⓘ	150.47	+4.67	+51.19
8.	8.	↓7.	Redis	Key-value, Multi-model ⓘ	146.53	+1.35	-7.58
9.	9.	9.	IBM Db2	Relational, Multi-model ⓘ	113.25	+1.87	-12.79
10.	↑11.	↑11.	Apache Cassandra	Wide column, Multi-model ⓘ	102.62	+0.74	-5.59
11.	↓10.	↓8.	Elasticsearch	Multi-model ⓘ	99.51	-4.07	-28.57
12.	12.	↓10.	SQLite	Relational	95.90	-0.08	-18.19
13.	13.	↑14.	MariaDB +	Relational, Multi-model ⓘ	86.41	-0.59	-7.93
14.	14.	↑16.	Microsoft Azure SQL Database	Relational, Multi-model ⓘ	75.96	+2.03	+0.59
15.	15.	↑18.	Apache Hive	Relational	74.13	+1.18	+1.48
16.	16.	↑17.	Splunk	Search engine	73.08	+0.87	-0.31
17.	17.	↓13.	Microsoft Access	Relational	64.92	-2.65	-30.45
18.	18.	↓15.	Amazon DynamoDB	Multi-model ⓘ	62.67	-2.75	-13.82
19.	19.	19.	Google BigQuery	Relational	52.36	-3.46	-8.37
20.	20.	20.	Neo4j	Graph	46.88	-0.44	-0.71
21.	↑22.	↑23.	SAP HANA	Relational, Multi-model ⓘ	33.23	+1.14	-0.17
22.	↓21.	↑24.	Apache Solr	Search engine, Multi-model ⓘ	33.17	+0.38	+0.77
23.	23.	↓22.	Teradata	Relational, Multi-model ⓘ	31.11	+0.03	-2.97
24.	24.	↓21.	FileMaker	Relational	28.81	-0.50	-7.67
25.	25.	25.	SAP Adaptive Server	Relational, Multi-model ⓘ	25.31	-0.22	-3.48
26.	↑28.	↑31.	ClickHouse	Relational, Multi-model ⓘ	22.87	+0.86	+4.14
27.	↓26.	↓26.	Microsoft Azure Cosmos DB	Multi-model ⓘ	22.43	-0.31	-0.12
28.	↓27.	↑29.	Apache Spark (SQL)	Relational	21.97	-0.10	+1.57
29.	29.	↑30.	PostGIS	Spatial, Multi-model ⓘ	21.51	-0.36	+1.90
30.	↑31.	↑32.	OpenSearch	Multi-model ⓘ	20.73	+0.46	+3.54

3) NOSQL DATABASES

NoSQL stands for:

- Not Just SQL
- No SQL Database

This type of database emerged from unmet needs. The internet was getting too big (too fast) for relational databases.

The name itself implies that there is no structure in the database... but is there?

But Why?

When we have a problem, we often try to use existing solutions. If they do not meet our needs, we try to develop something that does.

When specialized applications started demanding faster processing, NoSQL happened.

3) NOSQL DATABASES

Since it was developed to solve a specific problem, it had to bring some benefits to the table.

Larger web applications, such as online stores and social networks, had problems with the implementation and usage of relational databases.

Only 4 reasons were enough for NoSQL to exist and persist.

Four reasons why:

- **Scalability**
- **Cost**
- **Flexibility**
- **Availability**

Some features might be more important than others; depending on your application needs.

3) NOSQL DATABASES

Scalability:

Being able to meet the needs of various workloads is always a benefit. If you don't really need all the power all the time why pay for it?

For businesses/applications to grow they need to scale. There are two ways of scaling:

- Scaling Out
- Scaling Up


Scaling Out/Horizontal Scaling:

Adding resources per demand is called scaling out. Relational databases have trouble scaling out. Scaling out is adjusting to the current demand and then, going back.

Scaling Up/Vertical Scaling:

Is done by adding more RAM and CPU's. It also requires downtime since you have to make copies before you scale.

3) NOSQL DATABASES



Scaling Out/Horizontal Scaling:

Adding resources per demand is called scaling out. Relational databases have trouble scaling out. Scaling out is adjusting to the current demand and then, going back.

Scaling Up/Vertical Scaling:

Is done by adding more RAM and CPU's. It also requires downtime since you have to make copies before you scale.

3) NOSQL DATABASES

Cost:

When you use SQL solutions on such big scales the cost of software licensing increases.

There is also the cost of scaling, which is much easier done when scaling out vs scaling up.

NoSQL solutions are most often open source. Some companies offer open source software but charge for support, scaling or simply by usage.

How drastic is it?

On small applications the cost difference might be minimum but as the application grows costs grow as well.

Most of the cost is actually related to scalability but also to flexibility and availability.

3) NOSQL DATABASES

Flexibility:

Every DB administrator that worked with databases long enough experienced the headache of constraints, relations, tables.

We are, in essence, as flexible as our model supports us. You have probably learned this in Lab 3 when you start defining new tables and deleting old ones or change the datatype of an attribute.

Know your database...

With SQL you need to know your Database structure pretty much to work with it.

NoSQL – NoProblems :)

Some NoSQL solutions require no fixed table structures and are very flexible. You can add and remove as much as you want...

3) NOSQL DATABASES

Availability:

Time is money.

The more time your database is down the more money you waste.

With updates and maintenance, you often have to put your database down.

If your up-time is bad users tend to be hesitant to visit you again.

NoSQL is available!

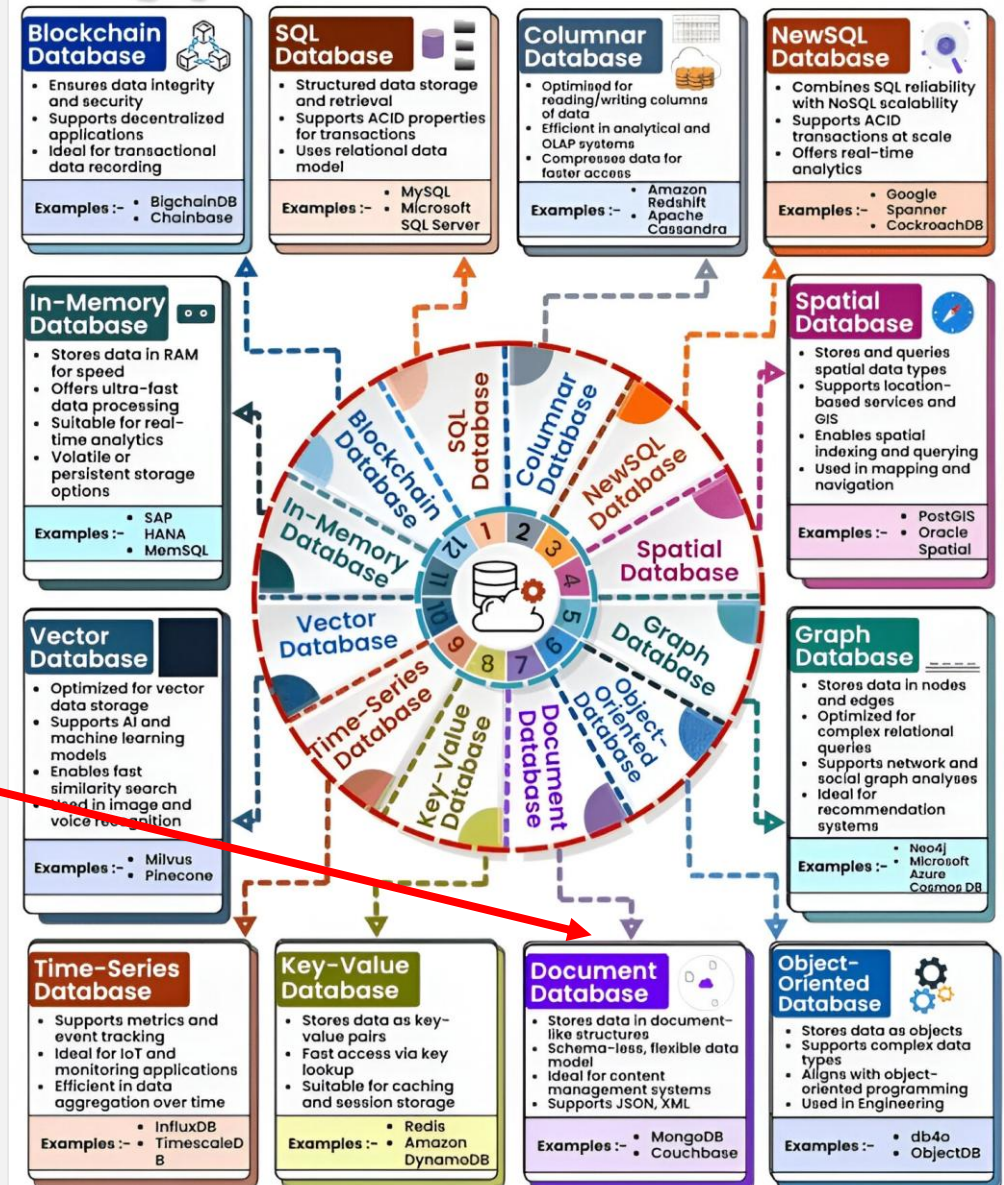
NoSQL databases can run on almost anything, they:

- have low-cost servers,
- work in clusters (one down, others take over),
- might suffer on performance but will be up!

5) GRAPH DATABASES

- There are many different types of databases, they can be sorted by type of data / type of application
- SQL Database, Graph Database, Document Database, Key-Value Database...

Types of Databases



4) DOCUMENT DATABASES

- A well-known (and well-ranked) document database is **MongoDB**. From MongoDB official site:

“**Document databases** offer a variety of advantages, including:

- An **intuitive data model** that is **fast** and **easy** for developers to work with
- A **flexible schema** that allows for the **data model to evolve** as application needs change
- The **ability** to horizontally **scale out**

Because of these advantages, document databases are general-purpose databases that can be used in a variety of use cases and industries.

Document databases are considered to be non-relational (or NoSQL) databases. Instead of storing data in fixed rows and columns, document databases use flexible documents. Document databases are the most popular alternative to tabular, relational databases.”

4) DOCUMENT DATABASES

- From MongoDB official site:

What are documents?

A document is a record in a document database. A document typically stores information about one object and any of its related metadata.

Documents store data in field-value pairs. The values can be a variety of types and structures, including strings, numbers, dates, arrays, or objects. Documents can be stored in formats like JSON, BSON, and XML.

4) DOCUMENT DATABASES

What is JSON?

JSON stands for JavaScript Object Notation but it is used in many more languages than Javascript because of its easy and powerful syntax.

To write JSON data, you **ONLY** need **7 characters**!

[...]: array

{ ... }: object

"key":"value": a key/value pair separated by :

" ... ": a character string

,: separator between objects and key/value pairs

All JSON is there:

[] { } : " ,

4) DOCUMENT DATABASES

What is BSON?

BSON stands for Binary JavaScript Object Notation (used by MongoDB): JSON is encoded as binary values:

JSON:

```
{
  "name": "Jerome Landre",
  "age": 40,
  "isStudent": false,
  "address": {
    "city": "Stockholm",
    "zip": "12345"
  },
  "hobbies": ["reading", "hiking", "coding"]
}
```

BSON:

```
5A 00 00 00 02 6E 61 6D 65 00 0D 00 00 00 4A 65 72 6F 6D 65 20 4C 61 6E 64 72 65 00 10 61
67 65 00 28 00 00 00 08 69 73 53 74 75 64 65 6E 74 00 00 03 61 64 64 72 65 73 73 00 2E 00
00 00 02 63 69 74 79 00 07 00 00 00 53 74 6F 63 6B 68 6F 6C 6D 00 02 7A 69 70 00 06 00 00
00 31 32 33 34 35 00 04 68 6F 62 62 69 65 73 00 1F 00 00 00 02 72 65 61 64 69 6E 67 00 08
00 00 00 02 68 69 6B 69 6E 67 00 08 00 00 00 02 63 6F 64 69 6E 67 00 00 00
```

4) DOCUMENT DATABASES

*	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	NUL	SOH	STX	ETX	EOT	ENQ	ACK	BEL	BS	TAB	LF	VT	FF	CR	SO	SI
1	DLE	DC1	DC2	DC3	DC4	NAK	SYN	ETB	CAN	EM	SUB	ESC	FS	GS	RS	US
2		!	"	#	\$	%	&	'	(	)	*	+	,	-	.	/
3	0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
4	@	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
5	P	Q	R	S	T	U	V	W	X	Y	Z	[	\	]	^	_
6	`	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
7	p	q	r	s	t	u	v	w	x	y	z	{		}	~	

JSON:

```
{
  "name": "Jerome Landre",
  "age": 40,
  "isStudent": false,
  "address": {
    "city": "Stockholm",
    "zip": "12345"
  },
  "hobbies": ["reading", "hiking", "coding"]
}
```

BSON:

```
5A 00 00 00 02 6E 61 6D 65 00 00 00 00 00 4A 65 72 6F 6D 65 20 4C 61 6E 64 72 65 00 10 61
67 65 00 28 00 00 00 08 69 73 53 74 75 64 65 6E 74 00 00 03 61 64 64 72 65 73 73 00 2E 00
00 00 02 63 69 74 79 00 07 00 00 00 53 74 6F 63 6B 68 6F 6C 6D 00 02 7A 69 70 00 06 00 00
00 31 32 33 34 35 00 04 68 6F 62 62 69 65 73 00 1F 00 00 00 02 72 65 61 64 69 6E 67 00 08
00 00 00 02 68 69 6B 69 6E 67 00 08 00 00 00 02 63 6F 64 69 6E 67 00 00 00
```

4) DOCUMENT DATABASES

Example of JSON data:

A JSON document that stores information about a user named Tom.

```
{  
  "id": 1,  
  "firstName": "Tom",  
  "lastName": "Sawyer",  
  "age": 17,  
  "gender": "male",  
  "address": {  
    "city": "St. Petersburg",  
    "state": "Missouri",  
    "country": "USA"  
  },  
  "hobbies": ["fishing", "adventures", "painting fences"],  
  "friends": ["Huckleberry Finn", "Becky Thatcher", "Joe Harper"]  
}
```

4) DOCUMENT DATABASES

Collections

A collection is a group of documents. Collections typically store documents that have similar contents.

Not all documents in a collection are required to have the same fields, because document databases have flexible schemas. Note that some document databases provide schema validation, so the schema can optionally be locked down when needed.

Continuing with the example above, the document with information about Tom could be stored in a collection named users. More documents could be added to the users collection in order to store information about other users. For example, the document below that stores information about Donna could be added to the users collection.

```
{
  "id": 1,
  "firstName": "Tom",
  "lastName": "Sawyer",
  "age": 17,
  "gender": "male",
  "address": {
    "city": "St. Petersburg",
    "state": "Missouri",
    "country": "USA"
  },
  "hobbies": ["fishing", "adventures", "painting fences"],
  "friends": ["Huckleberry Finn", "Becky Thatcher", "Joe Harper"]
}
```

```
{
  "id": 2,
  "firstName": "Rebecca",
  "lastName": "Thatcher",
  "nickName": "Becky",
  "age": 16,
  "gender": "female",
  "father": "Judge Thatcher",
  "school": "St. Petersburg Sunday School",
  "bestFriend": "Tom Sawyer",
  "personalityTraits": ["intelligent", "playful", "loyal"]
}
```

Collection “users”

```
[
  {
    "id": 1,
    "firstName": "Tom",
    "lastName": "Sawyer",
    "age": 17,
    "gender": "male",
    "address": {
      "city": "St. Petersburg",
      "state": "Missouri",
      "country": "USA"
    },
    "hobbies": ["fishing", "adventures", "painting fences"],
    "friends": ["Huckleberry Finn", "Becky Thatcher", "Joe Harper"]
  },
  {
    "id": 2,
    "firstName": "Rebecca",
    "lastName": "Thatcher",
    "nickName": "Becky",
    "age": 16,
    "gender": "female",
    "father": "Judge Thatcher",
    "school": "St. Petersburg Sunday School",
    "bestFriend": "Tom Sawyer",
    "personalityTraits": ["intelligent", "playful", "loyal"]
  }
]
```

4) DOCUMENT DATABASES

CRUD operations

Document databases typically have an API or query language that allows developers to execute the CRUD (create, read, update, and delete) operations.

- **Create:** Documents can be created in the database. Each document has a unique identifier.
- **Read:** Documents can be read from the database. The API or query language allows developers to query for documents using their unique identifiers or field values. Indexes can be added to the database in order to increase read performance.
- **Update:** Existing documents can be updated — either in whole or in part.
- **Delete:** Documents can be deleted from the database.collections

4) DOCUMENT DATABASES

What are the key features of document databases?

Document databases have the following key features:

- **Document model:** Data is stored in documents (unlike other databases that store data in structures like tables or graphs). Documents map to objects in most popular programming languages, which allows developers to rapidly develop their applications.
- **Flexible schema:** Document databases have flexible schemas, meaning that not all documents in a collection need to have the same fields. Note that some document databases support schema validation, so the schema can be optionally locked down.
- **Distributed and resilient:** Document databases are distributed, which allows for horizontal scaling (typically cheaper than vertical scaling) and data distribution. Document databases provide resiliency through replication.
- **Querying through an API or query language:** Document databases have an API or query language that allows developers to execute the CRUD operations on the database. Developers have the ability to query for documents based on unique identifiers or field values.

4) DOCUMENT DATABASES

MongoDB cheatsheet

Create Operations

Create or insert operations add new **documents** to a **collection**. If the collection does not currently exist, insert operations will create the collection.

MongoDB provides the following methods to insert documents into a collection:

- `db.collection.insertOne()`
- `db.collection.insertMany()`

In MongoDB, insert operations target a single **collection**. All write operations in MongoDB are **atomic** on the level of a single **document**.

```
db.users.insertOne(  ← collection
{
  name: "sue",        ← field: value
  age: 26,            ← field: value
  status: "pending"   ← field: value } document
})
```

4) DOCUMENT DATABASES

MongoDB cheatsheet

Read Operations

Read operations retrieve **documents** from a **collection**; i.e. query a collection for documents. MongoDB provides the following methods to read documents from a collection:

- `db.collection.find()`

You can specify **query filters or criteria** that identify the documents to return.

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

←	collection
←	query criteria
←	projection
←	cursor modifier

4) DOCUMENT DATABASES

MongoDB cheatsheet

Update Operations

Update operations modify existing **documents** in a **collection**. MongoDB provides the following methods to update documents of a collection:

- `db.collection.updateOne()`
- `db.collection.updateMany()`
- `db.collection.replaceOne()`

In MongoDB, update operations target a single collection. All write operations in MongoDB are **atomic** on the level of a single document.

You can specify criteria, or filters, that identify the documents to update. These **filters** use the same syntax as read operations.

```
db.users.updateMany(  
  { age: { $lt: 18 } },  
  { $set: { status: "reject" } }  
)
```

← collection
← update filter
← update action

4) DOCUMENT DATABASES

MongoDB cheatsheet

Delete Operations

Delete operations remove documents from a collection. MongoDB provides the following methods to delete documents of a collection:

- `db.collection.deleteOne()`
- `db.collection.deleteMany()`

In MongoDB, delete operations target a single **collection**. All write operations in MongoDB are **atomic** on the level of a single document.

You can specify criteria, or filters, that identify the documents to remove. These **filters** use the same syntax as read operations.

```
db.users.deleteMany(  
  { status: "reject" }  
)
```



4) DOCUMENT DATABASES

MongoDB playground

The screenshot shows the MongoDB Playground interface at <https://mongoplayground.net>. The interface is divided into three main sections: Database, Query, and Result.

- Database:** The database is set to 'bson'. It displays a list of documents. The first document (index 1565) has the following fields: `_id`: 93, `guid`: "3d0f18d3-0882-4877-a996-c...", `isActive`: false, `balance`: 2923.13, `picture`: "https://robohash.org/2...", `age`: 66, `eyeColor`: "green", `firstname`: "Nicole", `lastname`: "Patterson", `gender`: "female", `company`: "QUONATA", `email`: "nicolepatterson@quonata.", `phone`: "+1 (872) 460-3329", `address`: "456 Bridge Street, Bod...", and `registered`: "2018-12-06T05:32:32".
- Query:** The query entered is:


```
1 db.collection.find({
2   firstname: "Nicole"
3 })
```
- Result:** The result of the query is a single document:



```
{
  "_id": 93,
  "address": "456 Bridge Street, Bodega, Virgin",
  "age": 66,
  "balance": 2923.13,
  "company": "QUONATA",
  "email": "nicolepatterson@quonata.com",
  "eyeColor": "green",
  "firstname": "Nicole",
  "gender": "female",
  "guid": "3d0f18d3-0882-4877-a996-c274657737b!",
  "isActive": false,
  "lastname": "Patterson",
  "phone": "+1 (872) 460-3329",
  "picture": "https://robohash.org/25db5b4f-21...",
  "registered": "2018-12-06T05:32:32 -01:00"
}
```

At the bottom of the interface, it says "MongoDB version 6.0.20 - [Report an issue](#) - [About this playground](#)".

4) DOCUMENT DATABASES

MongoDB playground

To understand, please read the



MongoDB playground

The screenshot shows the MongoDB Playground web interface. At the top, there's a navigation bar with 'run', 'format', 'share', and 'docs' buttons. Below this, the interface is divided into three main sections: 'Database', 'Query', and 'Result'.

Database: A dropdown menu shows 'bson'. Below it, a list of documents is displayed, each with a unique '_id' and various fields like 'index', 'picture', 'age', 'fname', 'lname', 'gender', 'company', 'email', and 'phone'.

Query: A code editor shows a MongoDB query: `1 db.collection.find({ 2 age: { 3 $gt: 45 4 } 5 })`. The query is highlighted in blue.

Result: The result of the query is shown as a JSON array of documents. The first document is: `{ "_id": "69f21a9be3c918978109cbd1", "age": 49, "company": "ECRATIC", "email": "brittmooney@ecratic.com", "fname": "Britt", "gender": "male", "index": 1, "lname": "Mooney", "phone": "+1 (867) 579-3063", "picture": "http://placeholder.it/32x32" }`

At the bottom of the interface, there's a footer that reads: 'MongoDB version 8.0.21 - [Report an issue](#) - [About this playground](#)'.

Generate JSON data

json-generator.com/#

JSON GENERATOR Generate Reset Help Ukraine! Try the next version! Help Feedback

```
1 [
2   '{{repeat(5, 10)}}',
3   {
4     _id: '{{objectId()}}',
5     index: '{{index()}}',
6     picture: 'http://placeholder.it/32x32',
7     age: '{{integer(20, 50)}}',
8     fname: '{{firstName()}}',
9     lname: '{{surname()}}',
10    gender: '{{gender()}}',
11    company: '{{company().toUpperCase()}}',
12    email: '{{email()}}',
13    phone: '+1 {{phone()}}'
14  }
15 ]
```

1.10 KB

Copy JSON to clipboard

Download JSON file

Indent: 2 space tab

```
1 [
2   {
3     "_id": "69f21a9b091da372413d6601",
4     "index": 0,
5     "picture": "http://placeholder.it/32x32",
6     "age": 20,
7     "fname": "Blanchard",
8     "lname": "Herrera",
9     "gender": "male",
10    "company": "TETAK",
11    "email": "blanchardherrera@tetak.com",
12    "phone": "+1 (910) 448-3039"
13  },
14  {
15    "_id": "69f21a9be3c918978109cbb1",
16    "index": 1,
17    "picture": "http://placeholder.it/32x32",
18    "age": 49,
19    "fname": "Britt",
20    "lname": "Mooney",
21    "gender": "male",
22    "company": "ECRATIC",
23    "email": "brittmooney@ecratic.com",
24    "phone": "+1 (867) 579-3063"
25  },
26  {
27    "_id": "69f21a9b238eca8b3f7e944b",
28    "index": 2,
29    "picture": "http://placeholder.it/32x32",
30    "age": 45,
31    "fname": "Williams",
32    "lname": "Chan",
33    "gender": "male",
34    "company": "KYAGORO",
35    "email": "williamschan@kyagoro.com",
36    "phone": "+1 (873) 448-3836"
37  },
38  {
39    "_id": "69f21a9bed427554c59faa0a",
40    "index": 3,
41    "picture": "http://placeholder.it/32x32",
42    "age": 34,
43    "fname": "Ava",
44    "lname": "Curry",
45    "gender": "female",
46    "company": "ISOTEPNTA"
```


4) DOCUMENT DATABASES

**Let's go for a
live demonstration!**



In french:

Bonjour, ça va ?

Oui, ça va bien, merci !

Et toi ?

Ça va, merci.

/bɔ̃.ʒuʁ | sa va/

/wi | sa va bjɛ̃ | mɛʁ.si/

/e twa/

/sa va/

REMEMBER: BE POSITIVE!

- Mylène Farmer, “Désenchantée” (live version):

<https://www.youtube.com/watch?v=r1Le-dAocl8>

- Lyrics (English and Swedish translation):

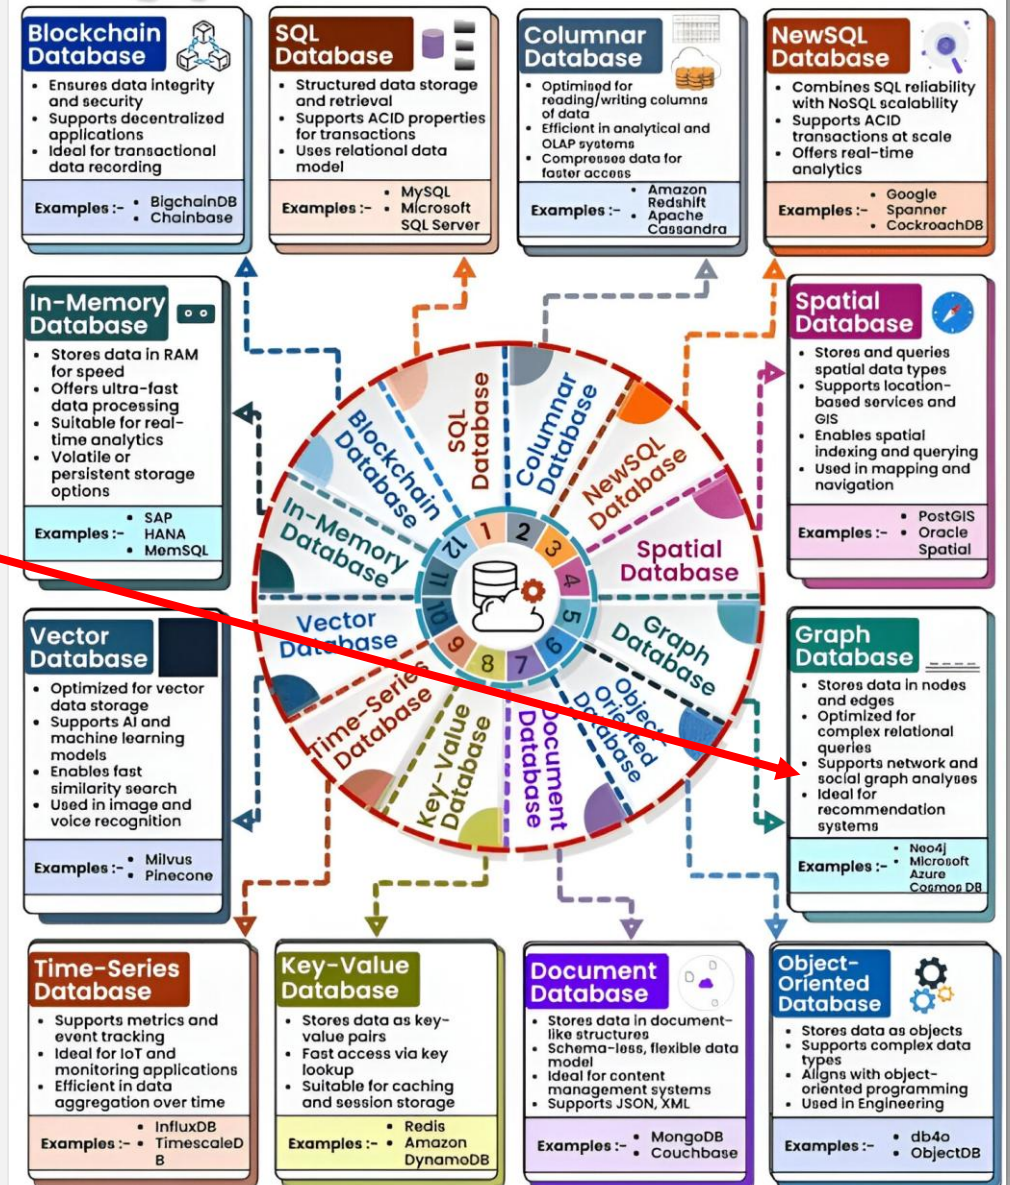
<https://lyricstranslate.com/en/desenchantee-disenchanted.html#songtranslation>

<https://lyricstranslate.com/en/desenchantee-besviken.html>

5) GRAPH DATABASES

- There are many different types of databases, they can be sorted by type of data / type of application
- SQL Database, Graph Database, Document Database, Key-Value Database

Types of Databases



6) NEO4J

The screenshot shows the Neo4j website homepage. At the top, there's a navigation bar with links for 'Aura Login', 'Partners', 'Company', 'Support', and a search icon. Below this is a secondary navigation bar with 'Products', 'Use Cases', 'Developers', 'AI Systems', 'Learn', and 'Pricing'. A 'Contact Us' link and a 'Get Started Free' button are also present.

The main content area features the Neo4j logo and the tagline 'THE WORLD'S LEADING GRAPH INTELLIGENCE PLATFORM'. Below this, the headline reads 'Build Intelligent Apps and AI For'. Underneath, it says 'AI-ready data by design' and provides two buttons: 'Start Building' and 'Learn More'.

A statistics section highlights three key metrics: '300k Developers Building', '80+ Fortune 100 Customers', and '170+ Partner Ecosystem'.


On the right side, there's a large grid of various graph analytics visualizations, including network graphs, bar charts, pie charts, and Cypher query snippets. Some of the queries shown are for finding fraud rings and analyzing relationships.

At the bottom, there's a row of logos for various partner companies, including Uber, Cisco, Walmart, Santander, Comcast, BNP Paribas, UBS, Airbus, BT Group, BMW, Boston Scientific, Novo Nordisk, and Dun & Bradstreet.

Source: <https://neo4j.com/>

6) NEO4J

Graph Database



Source: <https://neo4j.com/product/neo4j-graph-database>

The screenshot shows the Neo4j AuraDB landing page. At the top is a navigation bar with the Neo4j logo, links for Products, Use Cases, Developers, AI Systems, Learn, and Pricing, and buttons for Contact Us and Get Started Free. The main heading is 'Neo4j AuraDB: Fully Managed Graph Database'. Below it is a descriptive paragraph about mirroring data design and adapting to business needs. Two buttons, 'Start Free' and 'View Resources', are provided. To the right is a large blue graphic with icons representing graph data, a play button, and a server. Below this is a horizontal menu with links: How AuraDB Works, Capabilities, Use Cases, Customers, Plans, Partners, and Resources. The bottom section features four columns, each with an icon, a title, and a brief description of a key feature.

neo4j Products ▾ Use Cases ▾ Developers ▾ AI Systems Learn ▾ Pricing [Contact Us](#) [Get Started Free](#)

Neo4j AuraDB: Fully Managed Graph Database

Mirror your data design like sketching on a whiteboard and adapt effortlessly to changing business needs. Store data and relationships natively with a flexible schema.

[Start Free](#) [View Resources](#)

[How AuraDB Works](#) [Capabilities](#) [Use Cases](#) [Customers](#) [Plans](#) [Partners](#) [Resources](#)

Uncover Hidden Patterns
Reveal insights with Cypher queries and graph algorithms for your connected data.

Build Applications With Ease
Simplify development using a native graph model, flexible schema, and expressive Cypher.

Always-On, Zero Admin
99.95% uptime SLA, automated upgrades, patches, and maintenance.

Enterprise-Grade Security
Protect your data and ensure compliance with robust security features.





Get started for Free

Graph Intelligence for Contextual, AI-ready Data

- ✓ Fully-managed graph database, analytics, and agents
- ✓ Enterprise-grade security: SOC 2 Type II, ISO 27001 certified
- ✓ AI tools to help load and explore your data
- ✓ No credit card required

Also available via cloud marketplaces for simplified billing



Log in

Log in to Neo4j to continue

Log in

Don't have an account? [Sign up](#)

OR

 Continue with Google

 Continue with GitHub

 Continue with Organization SSO

 Continue with Microsoft

neo4j aura / 🏠 New Organization ▾ / 📁 New project ▾

🏠 Get started

</> Developer hub

Data services

🗄️ Instances

🔗 Graph Analytics

🔗 Data APIs

Tools

📁 Import

📄 Query

🔗 Explore

📊 Dashboards

📈 Operations >


📁 Project >

📖 Learning

Instances

Aura (0)

Self-managed (0)



Create your first instance

Configure your instance and start realizing the endless possibilities with graphs.

Create instance

Have data? [Import](#) existing data into Neo4j.

neo4j aura / New Organization / New project

Feedback

No instance connected Database: - User: -

Getting started

- Get started
- Developer hub
- Data services**
 - Instances
 - Graph Analytics
 - Data APIs
- Tools**
 - Import
 - Query
 - Explore
 - Dashboards
- Operations
- Project
- Learning**

Sample datasets

- Beginner
- Cypher reference
- Documentation
- Developer center

Learning pathways

- Graph databases
- Data analysis
- Reporting
- Software development
- Generative AI

Create an instance to run guide queries and import sample datasets.

Create instance

Sample datasets

GUIDE

Not Started

Social Network Analysis

- model Stack Overflow data as a graph;
- query the graph using Cypher;
- use shortest path algorithms.

GUIDE

Not Started

Crime investigation

- how crime data can be modeled;
- to refactor graph data;
- how to use spatial functions in Cypher.

GUIDE

Not Started

Healthcare analysis

- how Neo4j can support healthcare analytics;
- to query and filter data using Cypher;
- to aggregate data to find patterns.

GUIDE

Not Started

Movies

- create nodes and relationships;
- find nodes and patterns;
- understand and use the shortest path algorithm;
- write a basic recommendation query.

Connect to Neo4j

[Aura](#)[Self-managed](#)

**Connection to an instance is required to run queries or import data in this guide.**

Skipping this step will still allow you to view the guide.

[+ New connection](#)



No Aura instances found
Create your first Aura instance to get started

[Create instance](#)

[Continue](#)

neo4j aura / New Organization ▾ / New project ▾

Get started

Developer hub

Data services

Instances

Graph Analytics

Data APIs

Tools

Import

Query

Explore

Dashboards

Operations >

Project >

Learning

Create instance

AuraDB AuraDS

Confirm your tier. We'll recommend default values to help you get going. [See full feature comparison](#) ↗

☐ Professional

Free trial

Build production-ready applications

✓ Database monitoring with logs and metrics

✓ Predefined RBAC

✓ Backups, 7d retention

✓ Graph algorithms, vector optimized

✓ Standard support

From \$0.09 GB/hour

☐ Business Critical

Scale enterprise-grade applications

✓ Advanced monitoring, 3rd party integrations

✓ Custom RBAC, IP filtering, SSO

✓ Backups, 30d retention

✓ Graph algorithms, vector optimized

✓ 24/7 support

✓ High availability with 99.95% uptime SLA

From \$0.20 GB/hour

Min 2GB

☒ Free

Learn and explore graphs for free

✓ Up to 200k nodes and 400k relationships

✗ Limited memory and CPU

✗ Limited backups

✗ Auto-deleted after 30 days of inactivity

\$0 GB/hour

Deploy in a dedicated cloud environment — [See Virtual Dedicated Cloud](#) ↗

Instance details

Name

Instance01

-  Get started
-  Developer hub

Data services

Instances

Graph Analytics

Data APIs

Tools

Import

Query

Explore

Dashboards

Operations >

Project >

Learning

Instances

Aura (1)

Self-managed (0)

Search



Create instance

Instance01

● RUNNING

ID: be6b3a9f 



Connect ▾



Type: AuraDB Free Nodes: 0 Relationships: 0

Upgrade

▾ Metrics



neo4j aura / New Organization / New project

Feedback

Instance: Instance01 Database: be6b3a9f cypher 5 User: Aura (jerome.landre.info@gmail.com)

Get started
Developer hub
Data services
Instances
Graph Analytics
Data APIs
Tools
Import
Query
Explore
Dashboards
Operations
Project
Learning

Database information

Nodes (0)
Relationships (0)
Property keys

be6b3a9f\$

\$:welcome

GUIDE

Query fundamentals

You'll learn to:

- write basic queries;
- view graph and tabular results;
- perform queries to answer questions;
- perform advanced queries.

DATASET

Try Neo4j with the Movie Graph

Dive into a fully-built graph example to kickstart your Neo4j journey. Discover key query patterns through a fun, familiar world of actors and films.

Go to Learn Center

See the full catalogue of guides to help you get the most out of Neo4j.

[To Learn Center](#)

\$:connect

Connected to Instance01

neo4j auna / New Organization / New project

Query fundamentals

Step 2/7

Load the Northwind dataset

Again, if you already have completed the Data Importer guide, you can skip this step. If not, then first navigate to *Import*:

Go to Import

Use the button to load and map the data:

Load the Northwind dataset

Once that's all done, click Run import.

If you want to learn more about importing data, try the Importer guide.

With the data imported, navigate back to the *Query* tool and get started with learning Cypher:

Go back to Query

Get started

Developer hub

Data services

Instances

Graph Analytics

Data APIs

Tools

Import

Query

Explore

Dashboards

Operations

Project

Learning

Import

Data sources
Graph models
Import jobs



Connect your first data source

Choose from relational databases or CSV files to import your data into Neo4j.

New data source

neo4j aura / New Organization / New project

Feedback

Generate model

Run import

Graph models / Untitled model 1

Data source Local Browse

categories.csv

- categoryID 1
- categoryName Beverages
- description Soft drinks, cof...
- picture 0x151C2F0002...

customers.csv

employee-territories.csv

employees.csv

order-details.csv

orders.csv

products.csv

regions.csv

shippers.csv

suppliers.csv

territories.csv

Select a Node or Relationship

Query fundamentals

Step 2/7

Load the Northwind dataset

Again, if you already have completed the Data Importer guide, you can skip this step. If not, then first navigate to *Import*:

Go to Import

Use the button to load and map the data:

Load the Northwind dataset

Once that's all done, click **Run import**.

If you want to learn more about importing data, try the Importer guide.

With the data imported, navigate back to the Query tool and get started with learning Cypher:

Go back to Query

Get started

Developer hub

Data services

- Instances
- Graph Analytics
- Data APIs

Tools

- Import
- Query
- Explore
- Dashboards
- Operations
- Project
- Learning

Select instance



Select an instance to run your import

Instance

● Instance01

ID: be6b3a9f ▾

Cancel

Run import

Import results



Total time: 00:00:05



Run import completed

Import completed successfully.



Territory

territories.csv

[Show Cypher](#)

Time taken	File size	File rows	Nodes created	Properties set	Labels added	Query time
00:00:00	965 B	53	53	106	53	00:00:00



Region

regions.csv

[Show Cypher](#)

Time taken	File size	File rows	Nodes created	Properties set	Labels added	Query time
00:00:00	69 B	4	4	8	4	00:00:00



Supplier

suppliers.csv

[Show Cypher](#)

Time taken	File size	File rows	Nodes created	Properties set	Labels added	Query time
00:00:00	4.1 KiB	29	29	348	29	00:00:00


Close

Query fundamentals

Step 3/7

Get to know the database

A good place to start is to take stock of what you have in your database. You access *Database Information* from the top of the sidebar.



There you'll see how many nodes you have and all their labels, the number of relationships and their types, and also a list of all the property keys present in the database.

A node can have multiple labels and the list shows *all* labels in the database.

Relationships on the other hand, can have only one type, but multiple relationships can exist between the same nodes.

The list of property keys include both node properties and relationship properties.

You can click any of the labels, types, or property keys to run a query that returns a sample of nodes and/or relationships that has that label, type, or property.

For example, try click on the `Product` node label and you'll see this query populated and executed for you:

Previous

Next

 Get started Developer hub

Data services

 Instances Graph Analytics Data APIs

Tools

 Import Query Explore Dashboards Operations > Project > Learning

Graph models / Untitled model 3

Generate model

Show results



Run import



Data source Local Browse

categories.csv	...
categoryID	1
categoryName	Beverages
description	Soft drinks, cof...
picture	0x151C2F0002...

customers.csv ...

employee-territories.csv ...

employees.csv ...

order-details.csv ...

orders.csv ...


products.csv ...

regions.csv ...

shippers.csv ...

suppliers.csv ...

territories.csv ...



Select a Node or Relationship



neo4j aura / New Organization / New project

Feedback

Instance: Instance01 Database: be6b3a9f cypher 5 User: Aura (jerome.landre.info@gmail.com)

Database information

Nodes (1,104)

Category

Customer

Employee

Order

Product

Region

Shipper

Supplier

Territory

Relationships (4,909)

IN_REGION

IN_TERRITORY

ORDERS

PART_OF

PURCHASED

REPORTS_TO

SHIPS

SOLD

SUPPLIES

Property keys

address

birthDate

categoryID

categoryName

city

companyName

contactName

contactTitle

country

customerID

description

discontinued

discount

employeeID

extension

fax

firstName

freight

hireDate

homePage

Show all property keys (32 more)

1 MATCH path=(:Product)-[:PART_OF]->(c:Category)

2 WHERE c.categoryName = 'Produce'

3 RETURN path;

be6b3a9f\$ MATCH (n) RETURN n

No changes, no records

Completed after 7 ms

\$:welcome

GUIDE



Query fundamentals

You'll learn to:

- write basic queries;
- view graph and tabular results;
- perform queries to answer questions;
- perform advanced queries.

DATASET



Try Neo4j with the Movie Graph

Dive into a fully-built graph example to kickstart your Neo4j journey. Discover key query patterns through a fun, familiar world of actors and films.



Go to Learn Center

See the full catalogue of guides to help you get the most out of Neo4j.

Query fundamentals

Step 4/7

Excecuting queries with graph results

Central to the Query tool's UI is the Cypher editor on top where you write and run your queries and the list of result frames.

When you run any query, a result frame shows the query and its results, pushing previous ones down.

Try it out by typing in and executing the following Cypher query:

```
MATCH path=(:Product)-[:PART_OF]->(c:Category)
WHERE c.categoryName = 'Produce'
RETURN path;
```

You can edit the query within the frame and run it again without creating a new result frame.

The frame has two important views: Graph and Table. (Plan is for showing query plans and RAW for the raw results).

The default view is graph if your query returns graph elements like nodes, relationships or paths.

Relationships are only shown if you return a path, or name and return the relationships as well as the nodes.

Here you can see Product nodes with PART_OF relationships to the Category with name Produce.

On the right side of the graph view you can see the properties panel with information about the entities in your graph. You can click on any of them to change their styling, i.e. color, size, and caption, clicking on one element shows its attributes.

Previous

Next

Get started

Developer hub

Data services

Instances

Graph Analytics

Data APIs

Tools

Import

Query

Explore

Dashboards

Operations

Project

Learning

Last update: 10:51:22 AM

```
1 MATCH path=( :Product )-[ :PART_OF ]->(c:Category)
2 WHERE c.categoryName = 'Produce'
3 RETURN path;
```

The screenshot displays the Neo4j Aura web interface. On the left, a sidebar contains navigation links: 'Query fundamentals', 'Get started', 'Developer hub', 'Data services' (Instances, Graph Analytics, Data APIs), and 'Tools' (Import, Query, Explore, Dashboards, Operations, Project, Learning). The main workspace is divided into three panels. The top panel shows the Cypher editor with the query: `MATCH path=( :Product )-[ :PART_OF ]->(c:Category) WHERE c.categoryName = 'Produce' RETURN path;`. The middle panel, titled 'Database information', shows 1,104 nodes (Category, Customer, Employee, Order, Product, Region, Shipper, Supplier, Territory) and 4,909 relationships (IN_REGION, IN_TERRITORY, ORDERS, PART_OF, PURCHASED, REPORTS_TO, SHIPS, SOLD, SUPPLIES). It also lists property keys for the 'Produce' node. The bottom panel shows a graph visualization of the results, with a central 'Produce' node connected to five other nodes: 'Margim-up Drie...', 'Rösle Sauerkr...', 'Tofu', 'Uncle Bob's O...', and 'Longlife Tofu'. The right sidebar shows the 'Results overview' with 6 nodes and 5 relationships. The bottom status bar indicates 'Started streaming 5 records after 17 ms and completed after 17 ms.'

neo4jaura

New Organization

New project

Feedback

Query fundamentals

Step 5/7

Showing tabular results

If your query only returns scalar values (like strings or numbers), the result defaults to the table view and the graph view is not available.

Change your query to:

```
MATCH (p:Product)-[:PART_OF]->(c:Category)
WHERE c.categoryName = 'Produce'
RETURN p.productName, c.categoryName
```

The query uses variables, c and p, for the category and the product as you will want to refer to them later.

If you prefix your query with EXPLAIN or PROFILE you can see a plan option for a visual query plan, which helps later with optimizations. The RAW option shows the submitted query, the Neo4j server version and address, the result, and a summary.

Get started

Developer hub

Data services

Instances

Graph Analytics

Data APIs

Tools

Import

Query

Explore

Dashboards

Operations

Project

Learning

Instance: Instance01

Database: be6b3a9f CYPHER 5

User: Aura (jerome.landre.info@gmail.com)

Database information

Nodes (1,104)

Category

Customer

Employee

Order

Product

Region

Shipper

Supplier

Territory

Relationships (4,909)

IN_REGION

IN_TERRITORY

ORDERS

PART_OF

PURCHASED

REPORTS_TO

SHIPS

SOLD

SUPPLIES

Property keys

address

birthDate

categoryID

categoryName

city

companyName

contactName

contactTitle

country

customerID

description

discontinued

discount

employeeID

extension

fax

firstName

freight

hireDate

homePage

Show all property keys (32 more)

be6b3a9f\$

be6b3a9f\$ MATCH (p:Product)-[:PART_OF]->(c:Category) WHERE c.categoryName =

Table RAW

p.productName	c.categoryName
1 "Uncle Bob's Organic Dried Pears"	"Produce"
2 "Tofu"	"Produce"
3 "Rössle Sauerkraut"	"Produce"
4 "Manjimup Dried Apples"	"Produce"
5 "Longlife Tofu"	"Produce"

Started streaming 5 records after 17 ms and completed after 18 ms.

be6b3a9f\$ MATCH path=(:Product)-[:PART_OF]->(c:Category) WHERE c.categoryName

Graph Table RAW

Results overview

Nodes (6)

neo4jaura

New Organization

New project

Feedback

Instance: Instance01

Database: be6b3a9f cypher 5

User: Aura (jerome.landre.info@gmail.com)

Query fundamentals

Step 6/7

Use basic Cypher to answer questions

For writing correct queries it is always good to know what your data looks like. Besides the sidebar with the overview, a quick way to show the current graph schema is to CALL this procedure.

CALL db.schema.visualization()

Let's take a look at the different companies supplying Northwind with products and see which supplier provides products from which categories:

MATCH (s:Supplier)→(:Product)→c:Category
RETURN s.companyName as company, collect(c) as categories

This query finds suppliers that supply a product which is part of a category and returns the company names and the categories they supply products from, in a table. In cases where the same supplier supplies products from more than one category, the distinct category names are collected into a list, as you can see in the categories column.

Even though you specified the Product node in the

Get started

Developer hub

Data services

Instances

Graph Analytics

Data APIs

Tools

Import

Query

Explore

Dashboards

Operations

Project

Learning

Database information

Nodes (1,104)

Category Customer Employee Order Product Region Shipper Supplier Territory

Relationships (4,909)

IN_REGION IN_TERRITORY ORDERS PART_OF PURCHASED REPORTS_TO SHIPS SOLD SUPPLIES

Property keys

address birthDate categoryID categoryName city companyName contactName contactTitle country customerID description discontinued discount employeeID extension fax firstName freight hireDate homePage

Show all property keys (32 more)

be6b3a9f\$

be6b3a9f\$ CALL db.schema.visualization()

Graph Table RAW

```

graph TD
    Supplier -- supplies --> Product
    Product -- part of --> Order
    Order -- shipped to --> Customer
    Order -- shipped to --> Region
    Order -- shipped to --> Territory
    Customer -- reports to --> Region
    Region -- reports to --> Territory

```

Results overview

Nodes (9)


Category (1) Customer (1) Employee (1) Order (1) Product (1) Region (1) Shipper (1) Supplier (1) Territory (1)

Relationships (9)

IN_REGION (1) IN_TERRITORY (1) ORDERS (1) PART_OF (1) PURCHASED (1) REPORTS_TO (1) SHIPS (1) SOLD (1) SUPPLIES (1)

Started streaming 1 record after 2 ms and completed after 2 ms.

be6b3a9f\$ CALL db.schema.visualization🔗



Node details

Region

Key	Value	
<id>	-8540	🔗
constraints	["Constraint(id=5, name='regionID_Region_uniq', type='NO DE PROPERTY UNIQUENESS', schema=(:Region {regionID}), ownedIndex=4)"]	
indexes	[]	🔗
name	"Region"	🔗

nodes

[Customer , Category , Region , Supplier , Product , Shipper , Order , Territory , Employee]

relationships

[ORDERS , REPORTS_TO , PART_OF , IN_TERRITORY , IN_REGION , SUPPLIES , PURCHASED , SHIPS , SOLD]

Query fundamentals

Step 6/7

Use basic Cypher to answer questions

For writing correct queries it is always good to know what your data looks like. Besides the sidebar with the overview, a quick way to show the current graph schema is to `CALL` this procedure.

```
CALL db.schema.visualization()
```

Let's take a look at the different companies supplying Northwind with products and see which supplier provides products from which categories:

```
MATCH (s:Supplier)→(:Product)→c:Category
RETURN s.companyName as company, collect(c.categoryName) as categories
```

This query finds *suppliers* that *supply* a *product* which is *part of* a *category* and returns the company names and the categories they supply products from, in a table. In cases where the same supplier supplies products from more than one category, the distinct category names are collected into a list, as you can see in the *categories* column.

Even though you specified the *Product* node in the `MATCH` part of the query, no information about the actual product is returned since this is not included in the `RETURN` part.

Maybe you are interested only in suppliers of a certain category, say *Condiments*:

```
MATCH (s:Supplier)→(:Product)→c:Category
WHERE c.categoryName = 'Condiments'
RETURN DISTINCT s.companyName as company
```

- Get started
- Developer hub
- Data services
 - Instances
 - Graph Analytics
 - Data APIs
- Tools
 - Import
 - Query
 - Explore
 - Dashboards
 - Operations
 - Project
 - Learning

Database information

Nodes (1,104)

- Category
- Customer
- Employee
- Order
- Product
- Region
- Shipper
- Supplier
- Territory

Relationships (4,909)

- IN_REGION
- IN_TERRITORY
- ORDERS
- PART_OF
- PURCHASED
- REPORTS_TO
- SHIPS
- SOLD
- SUPPLIES

Property keys

- address
- birthDate
- categoryID
- categoryName
- city
- companyName
- contactName
- contactTitle
- country
- customerID
- description
- discontinued
- discount
- employeeID
- extension
- fax
- firstName
- freight
- hireDate
- homePage

Show all property keys (32 more)

Instance: Instance01 Database: be6b3a9f cypher 5 User: Aura (jerome.landre.info@gmail.com)

be6b3a9f\$

```
be6b3a9f$ MATCH (s:Supplier)-->(:Product)-->(c:Category) RETURN s.companyName as company
```

Table RAW

company	categories
1 "Exotic Liquids"	["Beverages", "Condiments"]
2 "New Orleans Cajun Delights"	["Condiments"]
3 "Grandma Kelly's Homestead"	["Condiments", "Produce"]
4 "Tokyo Traders"	["Meat/Poultry", "Seafood", "Produce"]
5 "Cooperativa de Quesos 'Las Cabras'"	["Dairy Products"]
6 "Mayumi's"	["Seafood", "Produce", "Condiments"]
7 "Pavlova"	["Confections", "Meat/Poultry", "Seafood", "Condiments", "Beverages"]
8 "Specialty Biscuits"	["Confections"]
9 "PB Knäckebröd AB"	["Grains/Cereals"]

Started streaming 29 records after 35 ms and completed after 36 ms.

```
be6b3a9f$ CALL db.schema.visualization()
```

Graph Table RAW

Node details

Region

Key	Value
<id>	-8540

Query fundamentals

Step 7/7

Write more advanced Cypher for problem-solving

Assume that you want to see which product categories are typically co-ordered with other product categories and how frequently.

This might help you understand which products to promote alongside others.

```
// which categories are the products of an order
MATCH (o:Order)-[:ORDERS]-(:Product)-[
// retain same ordering of categories
WITH o, c ORDER BY c.categoryName
// aggregate categories by order into a
WITH o, collect(DISTINCT c.categoryName
// only orders with more than one category
WHERE size(categories) > 1
// count how frequently the pairings occur
RETURN categories, count(*) as freq
// order by frequency
ORDER BY freq DESC
LIMIT 50
```

You can extend this and find customers who display similar purchasing patterns, i.e. which customers order similar products most frequently.

The base question is the same, you just expand across the product to other customers.

You find the "peer-groups" of our customers, which then can be used for product recommendations (people who bought X also bought) or segmentation into clusters of your customer base.

```
// pattern from customer purchasing products to another customer purchasing the
MATCH (c:Customer)-[:PURCHASED]-(:Order)-[
// don't want the same customer pair twice
```

Previous

Finish

Get started

Developer hub

Data services

Instances

Graph Analytics

Data APIs

Tools

Import

Query

Explore

Dashboards

Operations

Project

Learning

Instance: Instance01 Database: be6b3a9f cypher 5 User: Aura (jerome.landre.info@gmail.com)

Database information

Nodes (1,104)

Category Customer
Employee Order Product
Region Shipper Supplier
Territory

Relationships (4,909)

IN_REGION IN_TERRITORY
ORDERS PART_OF
PURCHASED REPORTS_TO
SHIPS SOLD SUPPLIES

Property keys

address birthDate categoryID
categoryName city
companyName contactName
contactTitle country
customerID description
discontinued discount
employeeID extension fax
firstName freight hireDate
homePage

Show all property keys (32 more)

be6b3a9f\$

be6b3a9f\$ // pattern from customer purchasing products to another customer purchasing the

Table RAW

	c.companyName	c2.companyName	similarity	topProducts
1	"Ernst Handel"	"Save-a-lot Markets"	184	["Guaraná Fantástica", "Gorgonzola Telino", "Alice Mutton", "Chang", "Pavlov a"]
2	"QUICK-Stop"	"Save-a-lot Markets"	150	["Chang", "Camembert Pierrot", "Singaporean Hokkien Fried Mee", "Konbu", "Gorgonzola Telino"]
3	"Rattlesnake Canyon Grocery"	"Save-a-lot Markets"	135	["Gnocchi di nonna Alice", "Alice Mutton", "Tarte au sucre", "Chang", "Gorgonzola Telino"]
4	"Ernst Handel"	"QUICK-Stop"	116	["Singaporean Hokkien Fried Mee", "Camembert Pierrot", "Wimmers gute Semmelknödel", "Gorgonzola Telino"]

Started streaming 10 records after 67 ms and completed after 188 ms.


be6b3a9f\$ // which categories are the products of an order in MATCH (o:Order)-[:ORDERS]-[

Table RAW

	categories	freq
1	["Beverages", "Confections"]	30
2	["Beverages", "Dairy Products"]	25
3	["Beverages", "Seafood"]	24
4	["Confections", "Dairy Products"]	23

7) CYPHER

Cypher is Neo4j's graph query language



The screenshot shows the Neo4j Cypher Guide interface. On the left is a dark sidebar with three icons: a database cylinder, a star, and a play button. The main content area has a dark background with white text. It starts with a back arrow icon and the title 'Cypher Guide'. Below this is the subtitle 'Neo4j's graph query language' and a paragraph stating 'Neo4j's Cypher language is purpose-built for working with graph data.' This is followed by a bulleted list of three features.

← **Cypher Guide**

Neo4j's graph query language

Neo4j's Cypher language is purpose-built for working with graph data.

- Uses patterns to describe graph data.
- Familiar SQL-like clauses.
- Declarative, describing what to find, not how to find it.

7) CYPHER

Cypher is Neo4j's graph query language

CREATE

Create a node

Let's use Cypher to generate a small social graph.

NOTE: This guide assumes that you use an empty graph.

1. Click this code block and bring it into the Editor:

```
CREATE (ee:Person {name: 'Emil', from: 'Sweden',  
kloutScore: 99})
```

- **CREATE** creates the node.
- **()** indicates the node.
- **ee:Person** – **ee** is the node variable and **Person** is the node label.
- **{}** contains the properties that describe the node.

2. Run the Cypher code by clicking the **Run** button.

7) CYPHER

Cypher is Neo4j's graph query language

CREATE more data

Nodes and relationships

The **CREATE** clause can create many nodes and relationships at once.

```
⌚ MATCH (ee:Person) WHERE ee.name = 'Emil'
  CREATE (js:Person { name: 'Johan', from: 'Sweden',
    learn: 'surfing' }),
    (ir:Person { name: 'Ian', from: 'England', title:
    'author' }),
    (rvb:Person { name: 'Rik', from: 'Belgium', pet:
    'Orval' }),
    (ally:Person { name: 'Allison', from:
    'California', hobby: 'surfing' }),
    (ee)-[:KNOWS {since: 2001}]->(js),(ee)-[:KNOWS
    {rating: 5}]->(ir),
```

7) CYPHER

Cypher is Neo4j's graph query language

MATCH patterns

Describe what to find in the graph

For instance, a pattern can be used to find Emil's friends:

```
⌕ MATCH (ee:Person)-[:KNOWS]-(friends)
  WHERE ee.name = 'Emil' RETURN ee, friends
```

- **MATCH** describes what nodes will be retrieved based upon the pattern.
- **(ee)** is the node reference that will be returned based upon the **WHERE** clause.
- **-[:KNOWS]-** matches the **KNOWS** relationships (in either direction) from **ee**.
- **(friends)** represents the nodes that are Emil's friends.
- **RETURN** returns the node, referenced here by **(ee)**, and the related **(friends)** nodes found.

7) CYPHER

Cypher is Neo4j's graph query language

Next steps

- [🔗 The Movie Graph](#) – Queries and recommendations with Cypher - movie use case.
- [🔗 The Northwind Graph](#) – Translate and import relation data into graph.

References


- [🔗 Help commands](#) – Useful Neo4j Browser commands
- [🔗 Help keys](#) – Keyboard shortcuts

7) CYPHER

Cypher contains many instructions

In this course, we will focus mainly on:

- **MATCH** (the Read in the CRUD, equivalent to SELECT)
- **MERGE** (the Create in the CRUD, equivalent to INSERT)



The screenshot shows the Neo4j Cypher Cheat Sheet for Version 5. The left sidebar lists various query types: Read Query (selected), Write query, Patterns, and Clauses. The main content area is titled 'Read Query' and includes a 'Read Query Structure' section with a list of clauses: [USE], [MATCH [WHERE]], [OPTIONAL MATCH [WHERE]], [WITH [ORDER BY] [SKIP] [LIMIT] [WHERE]], and RETURN [ORDER BY] [SKIP] [LIMIT]. Below this, there are examples for the MATCH clause:

- MATCH**: Match all nodes and return all nodes. Example: `MATCH (n) RETURN n`
- MATCH**: Find all nodes with the `Movie` label. Example: `MATCH (movie:Movie) RETURN movie.title`
- MATCH**: Find the types of an aliased relationship. Example: `MATCH (:Person {name: 'Oliver Stone'})-[r]->() RETURN type(r) AS relType`
- MATCH**: Relationship pattern filtering on the `ACTED_IN` relationship type. Example: `MATCH (:Movie {title: 'Wall Street'})<-[:ACTED_IN]-(actor:Person) RETURN actor.name AS actor`

ANY QUESTIONS?

BE POSITIVE, B+!